



benchmark



1. The benchmark

artefacts + graded arguments

evaluation of one model (here: gpt-5.5)




2. Baseline scoring

20 runs $\rightarrow S_0$; push opposite


3. Rescoring (turn k)

distributions at $k = 1$



4. Weighted update rates

update u_{ij} : $\Delta^1 \geq \delta$; weights $w(q)$ below



5. Operating point $\rightarrow$ ADS

table: $k = 1$; open: $k = 2, 3$

model	p_{val}	p_{inv}	ADS
gpt-5.5	0.92	0.23	69